

```
In [1]: import os
os.environ["PYSPARK_PYTHON"] = "/home/shiva/venv/bin/python"
os.environ["PYSPARK_DRIVER_PYTHON"] = "/home/shiva/venv/bin/python"
```

# Practice Questions Solutions

## Sources

- Questions source - ChatGPT, (Questions.md)
- Official answer source - ChatGPT, (Answers.md)
- CSV source - Kaggle - <https://www.kaggle.com/datasets/urvishahir/electric-vehicle-specifications-dataset-2025>

## Setup

```
In [51]: import pyspark.sql as ps
import pyspark.sql.functions as psf
import pyspark.sql.types as pst
import pandas as pd
import pyspark.sql.window as psw
```

```
In [3]: spark = (ps.SparkSession
                .builder
                .appName("EV Spec Practice")
                .getOrCreate()
            )
```

```
WARNING: Using incubator modules: jdk.incubator.vector
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
25/07/03 22:55:17 WARN Utils: Your hostname, Victus, resolves to a loopback address: 127.0.1.1; using 192.168.29.87 instead (on interface wlo1)
25/07/03 22:55:17 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
25/07/03 22:55:18 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

```
In [4]: file_path = "electric_vehicles_spec_2025.csv.csv"
```

```
In [5]: df = (
    spark
    .read
    .option("samplingRatio", 0.001)
    .option("header", "true")
    .option("inferSchema", "true")
    .csv(file_path)
)
```

```
In [6]: df.columns
```

```
Out[6]: ['brand',
         'model',
         'top_speed_kmh',
         'battery_capacity_kWh',
         'battery_type',
         'number_of_cells',
         'torque_nm',
         'efficiency_wh_per_km',
         'range_km',
         'acceleration_0_100_s',
         'fast_charging_power_kw_dc',
         'fast_charge_port',
         'towing_capacity_kg',
         'cargo_volume_l',
         'seats',
         'drivetrain',
         'segment',
         'length_mm',
         'width_mm',
         'height_mm',
         'car_body_type',
         'source_url']
```

```
In [7]: (
         df
         .createOrReplaceTempView("df_view")
         )
```

---

## Easy Questions

Focus: RDD/DataFrame basics, filtering, selecting, grouping, sorting, basic aggregations, handling nulls.

### DataFrame Basics

1. Load the CSV file into a DataFrame and print the schema.

```
df = (
    spark
    .read
    .option("samplingRatio", 0.001)
    .option("header", "true")
    .option("inferSchema", "true")
    .csv(file_path)
)````
```

```
In [52]: df.printSchema()
```

```
root
|-- brand: string (nullable = true)
|-- model: string (nullable = true)
|-- top_speed_kmh: integer (nullable = true)
|-- battery_capacity_kWh: double (nullable = true)
|-- battery_type: string (nullable = true)
|-- number_of_cells: integer (nullable = true)
|-- torque_nm: integer (nullable = true)
|-- efficiency_wh_per_km: integer (nullable = true)
|-- range_km: integer (nullable = true)
|-- acceleration_0_100_s: double (nullable = true)
|-- fast_charging_power_kw_dc: integer (nullable = true)
|-- fast_charge_port: string (nullable = true)
|-- towing_capacity_kg: integer (nullable = true)
|-- cargo_volume_l: integer (nullable = true)
|-- seats: integer (nullable = true)
|-- drivetrain: string (nullable = true)
|-- segment: string (nullable = true)
|-- length_mm: integer (nullable = true)
|-- width_mm: integer (nullable = true)
|-- height_mm: integer (nullable = true)
|-- car_body_type: string (nullable = true)
|-- source_url: string (nullable = true)
```

```
In [53]: df.schema
```

```
Out[53]: StructType([StructField('brand', StringType(), True), StructField('model', StringType(), True), StructField('top_speed_kmh', IntegerType(), True), StructField('battery_capacity_kWh', DoubleType(), True), StructField('battery_type', StringType(), True), StructField('number_of_cells', IntegerType(), True), StructField('torque_nm', IntegerType(), True), StructField('efficiency_wh_per_km', IntegerType(), True), StructField('range_km', IntegerType(), True), StructField('acceleration_0_100_s', DoubleType(), True), StructField('fast_charging_power_kw_dc', IntegerType(), True), StructField('fast_charge_port', StringType(), True), StructField('towing_capacity_kg', IntegerType(), True), StructField('cargo_volume_l', IntegerType(), True), StructField('seats', IntegerType(), True), StructField('drivetrain', StringType(), True), StructField('segment', StringType(), True), StructField('length_mm', IntegerType(), True), StructField('width_mm', IntegerType(), True), StructField('height_mm', IntegerType(), True), StructField('car_body_type', StringType(), True), StructField('source_url', StringType(), True)])
```

---

2. Display the first 10 rows of the DataFrame.

```
In [54]: df.show(10)
```

```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
| brand|                model|top_speed_kmh|battery_capacity_kWh|battery_ty
pe|number_of_cells|torque_nm|efficiency_wh_per_km|range_km|acceleration_0_
100_s|fast_charging_power_kw_dc|fast_charge_port|towing_capacity_kg|cargo_
volume_l|seats|drivetrain|      segment|length_mm|width_mm|height_mm|car_bo
dy_type|                source_url|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Abarth|    500e Convertible|    155|    37.8| Lithium-i
on|    192|    235|    156|    225|
7.0|    67|    CCS|    0|
185|    4|    FWD| B - Compact|    3673|    1683|    1518|    Hatchba
ck|https://ev-databa...|
|Abarth|    500e Hatchback|    155|    37.8| Lithium-i
on|    192|    235|    149|    225|
7.0|    67|    CCS|    0|
185|    4|    FWD| B - Compact|    3673|    1683|    1518|    Hatchba
ck|https://ev-databa...|
|Abarth| 600e Scorpionissima|    200|    50.8| Lithium-i
on|    102|    345|    158|    280|
5.9|    79|    CCS|    0|
360|    5|    FWD|JB - Compact|    4187|    1779|    1557|    S
UV|https://ev-databa...|
|Abarth|    600e Turismo|    200|    50.8| Lithium-i
on|    102|    345|    158|    280|
6.2|    79|    CCS|    0|
360|    5|    FWD|JB - Compact|    4187|    1779|    1557|    S
UV|https://ev-databa...|
|Aiways|    U5|    150|    60.0| Lithium-i
on|    NULL|    310|    156|    315|
7.5|    78|    CCS|    NULL|
496|    5|    FWD| JC - Medium|    4680|    1865|    1700|    S
UV|https://ev-databa...|
|Aiways|    U6|    160|    60.0| Lithium-i
on|    NULL|    315|    150|    350|
7.0|    78|    CCS|    NULL|
472|    5|    FWD| JC - Medium|    4805|    1880|    1641|    S
UV|https://ev-databa...|
| Alfa|Romeo Junior Elet...|    150|    50.8| Lithium-i
on|    102|    260|    128|    320|
9.0|    85|    CCS|    0|
400|    5|    FWD|JB - Compact|    4173|    1781|    1532|    S
UV|https://ev-databa...|
| Alfa|Romeo Junior Elet...|    200|    50.8| Lithium-i
on|    102|    345|    164|    310|
6.0|    85|    CCS|    0|
400|    5|    FWD|JB - Compact|    4173|    1781|    1505|    S
UV|https://ev-databa...|
|Alpine|A290 Electric 180 hp|    160|    52.0| Lithium-i
on|    184|    285|    138|    310|
7.4|    70|    CCS|    500|
326|    5|    FWD| B - Compact|    3997|    1823|    1512|    Hatchba
ck|https://ev-databa...|

```

Alpine	A290 Electric	220 hp	170	52.0	Lithium-i
on	184	300	144	305	
6.4		70	CCS	500	
326	5	FWD	B - Compact	3997	1823
1512					Hatchba

ck|https://ev-databa...

```

+-----+-----+-----+-----+-----+
--+-+-----+-----+-----+-----+-----+
-----+-+-----+-----+-----+-----+
-----+-+-----+-----+-----+-----+
-----+-+-----+-----+-----+-----+
-----+-+-----+-----+-----+-----+

```

only showing top 10 rows

3. Select only `brand`, `model`, and `range_km` columns.

```
In [55]: ans_df_3 = (
          df
          .select("brand", "model", "range_km")
        )
```

```
In [56]: ans_df_3.show(10)
```

brand	model	range_km
Abarth	500e Convertible	225
Abarth	500e Hatchback	225
Abarth	600e Scorpionissima	280
Abarth	600e Turismo	280
Aiways	U5	315
Aiways	U6	350
Alfa Romeo	Junior Elet...	320
Alfa Romeo	Junior Elet...	310
Alpine	A290 Electric 180 hp	310
Alpine	A290 Electric 220 hp	305

only showing top 10 rows

4. Rename the column `battery_capacity_kWh` to `battery_kWh`.

```
In [58]: df_renamed = (
          df
          .withColumnRenamed("battery_capacity_kwh", "battery_kwh")
        )
```

```
In [61]: df_renamed.columns
```

```
Out[61]: ['brand',
          'model',
          'top_speed_kmh',
          'battery_kwh',
          'battery_type',
          'number_of_cells',
          'torque_nm',
          'efficiency_wh_per_kmh',
          'range_kmh',
          'acceleration_0_100_s',
          'fast_charging_power_kw_dc',
          'fast_charge_port',
          'towing_capacity_kg',
          'cargo_volume_l',
          'seats',
          'drivetrain',
          'segment',
          'length_mm',
          'width_mm',
          'height_mm',
          'car_body_type',
          'source_url']
```

---

5. Drop the `source_url` column from the DataFrame.

```
In [62]: dropped_source_url_df = (
          df
          .drop("source_url")
          )
```

```
In [66]: len(df.columns)
```

```
Out[66]: 22
```

```
In [67]: len(dropped_source_url_df.columns)
```

```
Out[67]: 21
```

---

## Filtering & Selection

6. Filter rows where `brand` is "Abarth".

```
In [72]: (
          df
          .filter(psf.expr("brand == 'Abarth'"))
          .select("brand", "model")
          .show(10)
          )
```

brand	model
Abarth	500e Convertible
Abarth	500e Hatchback
Abarth	600e Scorpionissima
Abarth	600e Turismo

7. Find all vehicles with `top_speed_kmh` greater than 180.

```
In [79]: query = """
        SELECT
            brand,
            model,
            top_speed_kmh
        FROM
            df_view
        WHERE
            top_speed_kmh > 180
        """
```

```
In [80]: spark.sql(query).show(10)
```

brand	model	top_speed_kmh
Abarth	600e Scorpionissima	200
Abarth	600e Turismo	200
Alfa Romeo	Junior Elet...	200
Audi	A6 Avant e-tron	210
Audi	A6 Avant e-tron p...	210
Audi	A6 Avant e-tron q...	210
Audi	A6 Sportback e-tron	210
Audi	A6 Sportback e-tr...	210
Audi	A6 Sportback e-tr...	210
Audi	Q6 e-tron	210

only showing top 10 rows

8. Filter rows where `torque_nm` is not null.

```
In [83]: (
    df
    .select("brand", "model", "torque_nm")
    .filter(psf.col("torque_nm").isNotNull())
    .show(5)
)
```

brand	model	torque_nm
Abarth	500e Convertible	235
Abarth	500e Hatchback	235
Abarth	600e Scorpionissima	345
Abarth	600e Turismo	345
Aiways	U5	310

only showing top 5 rows

9. Retrieve vehicles where `range_km` is between 250 and 300.

```
In [85]: query = """
        SELECT
            brand,
            model,
            range_km
        FROM
            df_view
        WHERE
            range_km BETWEEN 250 AND 300
        ORDER BY
            range_km
        """
```

```
In [86]: spark.sql(query).show(10)
```

brand	model	range_km
Mini	Cooper E	250
firefly	NULL	250
Renault	5 E-Tech 40kWh 120hp	250
Smart	#1 Pro	250
Smart	#1 Pure	250
Dongfeng	Box 42.3 kWh	255
Citroen	e-C3	255
Hyundai	INSTER Standard R...	255
Renault	5 E-Tech 40kWh 95hp	255
Citroen	e-SpaceTourer M 7...	260

only showing top 10 rows

10. Filter SUVs with more than 4 seats.

```
In [93]: (
    df
    .select("brand", "model", "car_body_type", "seats")
    .filter(psf.col("car_body_type") == "SUV")
    .filter(psf.col("seats") > 4)
```



```

        .orderBy("seats")
        .show(10)
    )

```

```

+-----+-----+-----+-----+
| brand|          model|car_body_type|seats|
+-----+-----+-----+-----+
|Abarth| 600e Scorpionissima|      SUV|    5|
|Abarth|      600e Turismo|      SUV|    5|
|Aiways|              U5|      SUV|    5|
|Aiways|              U6|      SUV|    5|
|  Alfa|Romeo Junior Elet...|      SUV|    5|
|  Alfa|Romeo Junior Elet...|      SUV|    5|
| Audi|Q4 Sportback e-tr...|      SUV|    5|
| Audi|Q4 Sportback e-tr...|      SUV|    5|
| Audi|Q4 Sportback e-tr...|      SUV|    5|
| Audi|Q4 Sportback e-tr...|      SUV|    5|
+-----+-----+-----+-----+
only showing top 10 rows

```

---

## Aggregations

11. Calculate the average `efficiency_wh_per_km` for each `brand` .

```

In [99]: (
    df
    .groupBy("brand")
    .avg("efficiency_wh_per_km")
    .show(10)
)

```

```

+-----+-----+
|      brand|avg(efficiency_wh_per_km)|
+-----+-----+
| Leapmotor|             151.0|
| Volkswagen| 150.08695652173913|
| Peugeot| 159.3684210526316|
|  NI0| 181.08333333333334|
|  Lexus| 150.66666666666666|
| Polestar| 156.44444444444446|
|  Jaguar|             208.0|
| Maserati|             200.0|
|Rolls-Royce|             192.0|
|   Jeep|             140.0|
+-----+-----+
only showing top 10 rows

```

```

In [105... (
    df
    .groupBy("brand")
    .agg(psf.avg("efficiency_wh_per_km").alias("average"))
    .show(10)
)

```

```

+-----+-----+
|      brand|      average|
+-----+-----+
|  Leapmotor|      151.0|
| Volkswagen|150.08695652173913|
|   Peugeot| 159.3684210526316|
|        NIO|181.08333333333334|
|    Lexus|150.66666666666666|
| Polestar|156.44444444444446|
|   Jaguar|      208.0|
| Maserati|      200.0|
|Rolls-Royce|      192.0|
|    Jeep|      140.0|
+-----+-----+
only showing top 10 rows

```

---

12. Find the maximum `top_speed_kmh` for each `brand`.

```

In [106... (
    df
    .select("brand", "top_speed_kmh")
    .groupBy("brand")
    .agg(psf.max("top_speed_kmh").alias("Max_speed"))
    .show(10)
)

```

```

+-----+-----+
|      brand|Max_speed|
+-----+-----+
|  Leapmotor|      170|
| Volkswagen|      200|
|   Peugeot|      170|
|        NIO|      200|
|    Lexus|      160|
| Polestar|      210|
|   Jaguar|      200|
| Maserati|      325|
|Rolls-Royce|      250|
|    Jeep|      180|
+-----+-----+
only showing top 10 rows

```

---

13. Count how many distinct `car_body_type` values exist.

```

In [107... (
    df
    .agg(psf.count_distinct("car_body_type").alias("distinct num"))
    .show()
)

```

```

+-----+
|distinct num|
+-----+
|          8|
+-----+

```

```

In [108... (
    df
    .select("car_body_type")
    .distinct()
    .count()
)

```

Out[108... 8

---

14. Compute the total number of vehicles for each `segment` .

```

In [113... (
    df
    .groupBy("segment")
    .count()
    .show(10)
)

```

```

+-----+-----+
|          segment|count|
+-----+-----+
|    JB - Compact|   44|
|N - Passenger Van|   47|
|      A - Mini|    3|
|    B - Compact|   29|
|    F - Luxury|   51|
|    JD - Large|   58|
|    C - Medium|   34|
|    D - Large|   28|
|    E - Executive|  30|
|    JF - Luxury|   30|
+-----+-----+
only showing top 10 rows

```

---

15. Group by `drivetrain` and calculate the average `acceleration_0_100_s` .

```

In [115... (
    df
    .groupBy("drivetrain")
    .agg(psf.avg("acceleration_0_100_s").alias("average"))
    .show(10)
)

```

```

+-----+-----+
|drivetrain|          average|
+-----+-----+
|      FWD| 9.730128205128207|
|      AWD|4.5539267015706795|
|      RWD| 6.887022900763362|
+-----+-----+

```

---

## Sorting

16. Sort the DataFrame by `range_km` in descending order.

```

In [117... (
    df
      .select("brand", "model", "range_km")
      .sort("range_km", ascending=False)
      .show(10)
)

```

```

+-----+-----+-----+
|      brand|      model|range_km|
+-----+-----+-----+
|Mercedes-Benz|      EQS 450+|      685|
|      Lucid|  Air Grand Touring|      665|
|Mercedes-Benz|      EQS 450 4MATIC|      655|
|Mercedes-Benz|      EQS 500 4MATIC|      640|
|Mercedes-Benz|      EQS 580 4MATIC|      640|
|      Audi|A6 Sportback e-tr...|      610|
|      Audi|A6 Sportback e-tr...|      590|
|Mercedes-Benz|  EQS AMG 53 4MATIC+|      585|
|      Lucid|      Air Touring|      580|
|      Audi|A6 Avant e-tron p...|      575|
+-----+-----+-----+
only showing top 10 rows

```

---

17. Sort vehicles by `battery_capacity_kWh`, then by `top_speed_kmh`.

```

In [121... (
    df
      .select("brand", "model", "battery_capacity_kWh", "top_speed_kmh")
      .orderBy("battery_capacity_kWh", "top_speed_kmh").show()
)

```

brand	model	battery_capacity_kWh	top_speed_kmh
Fiat	500e 3+1 24 kWh	21.3	135
Fiat	500e Cabrio 24 kWh	21.3	135
Fiat	500e Hatchback 24...	21.3	135
Dacia	Spring Electric 65	25.0	125
Dacia	Spring Electric 45	25.0	125
Dongfeng	Box 31.4 kWh	29.0	140
BYD	DOLPHIN SURF 30 k...	30.0	150
Leapmotor	T03	36.0	130
Mini	Cooper E	36.6	160
Fiat	500e 3+1 42 kWh	37.3	150
Fiat	500e Cabrio 42 kWh	37.3	150
Fiat	500e Hatchback 42...	37.3	150
Abarth	500e Convertible	37.8	155
Abarth	500e Hatchback	37.8	155
Mini	Aceman E	38.5	160
Hyundai	INSTER Standard R...	39.0	140
Renault	5 E-Tech 40kWh 95hp	40.0	130
Dongfeng	Box 42.3 kWh	40.0	140
Renault	4 E-Tech 40kWh 120hp	40.0	150
Renault	5 E-Tech 40kWh 120hp	40.0	150

only showing top 20 rows

## Handling Nulls

18. Count how many null values are present in each column.

```
In [11]: (
    df
    .select([psf.sum(psf.col(c).isNull()).cast("int")).alias(c) for c in d
    .transpose()
    .show(22)
)
```

key	0
model	1
top_speed_kmh	0
battery_capacity_kWh	0
battery_type	0
number_of_cells	202
torque_nm	7
efficiency_wh_per_kmh	0
range_kmh	0
acceleration_0_100_s	0
fast_charging_power_kw	1
fast_charge_port	1
towing_capacity_kg	26
cargo_volume_l	4
seats	0
drivetrain	0
segment	0
length_mm	0
width_mm	0
height_mm	0
car_body_type	0
source_url	0

---

19. Drop rows with any null values.

```
In [127... df_dropped_rows = (
    df
    .na
    .drop()
)
```

```
In [130... (
    df_dropped_rows
    .select("brand", "model")
    .show(10)
)
```

```

+-----+-----+
| brand|          model|
+-----+-----+
|Abarth|    500e Convertible|
|Abarth|      500e Hatchback|
|Abarth|  600e Scorpionissima|
|Abarth|      600e Turismo|
|  Alfa|Romeo Junior Elet...|
|  Alfa|Romeo Junior Elet...|
|Alpine|A290 Electric 180 hp|
|Alpine|A290 Electric 220 hp|
|  Audi|      A6 Avant e-tron|
|  Audi|A6 Avant e-tron p...|
+-----+-----+
only showing top 10 rows

```

---

20. Fill nulls in `number_of_cells` with 0.

```

In [131... df_filled = (
    df
    .select("number_of_cells")
    .fillna(0)
)

```

```

In [132... (
    df_filled
    .select("number_of_cells")
    .show(10)
)

```

```

+-----+
|number_of_cells|
+-----+
|          192|
|          192|
|          102|
|          102|
|             0|
|             0|
|          102|
|          102|
|          184|
|          184|
+-----+
only showing top 10 rows

```

---

21. Fill nulls in `towing_capacity_kg` with the mean value of the column.

```

In [140... mean_val = (
    df
    .select(psf.mean("towing_capacity_kg"))
)

```

```
.first()
)[0]
```

```
In [141]: df_filled_with_mean = (
            df
            .select("brand", "model", "towing_capacity_kg")
            .fillna({"towing_capacity_kg": mean_val})
        )
```

```
In [142]: df_filled_with_mean.show(10)
```

```
+-----+-----+-----+
| brand|          model|towing_capacity_kg|
+-----+-----+-----+
|Abarth|   500e Convertible|          0|
|Abarth|   500e Hatchback|          0|
|Abarth| 600e Scorpionissima|          0|
|Abarth|   600e Turismo|          0|
|Aiways|          U5|        1052|
|Aiways|          U6|        1052|
| Alfa|Romeo Junior Elet...|          0|
| Alfa|Romeo Junior Elet...|          0|
|Alpine|A290 Electric 180 hp|          500|
|Alpine|A290 Electric 220 hp|          500|
+-----+-----+-----+
only showing top 10 rows
```

---

## Distinct & Deduplication

22. Get distinct values from the `brand` column.

```
In [10]: (
            df
            .select("brand")
            .distinct()
            .show(10)
        )
```

```
+-----+
|      brand|
+-----+
| Leapmotor|
| Volkswagen|
| Peugeot|
|      NIO|
|      Lexus|
| Polestar|
|      Jaguar|
| Maserati|
|Rolls-Royce|
|      Jeep|
+-----+
only showing top 10 rows
```

---



23. Drop duplicate rows from the DataFrame.

```
In [12]: df_without_duplicates = (  
    df  
    .dropDuplicates()  
    )
```

```
In [14]: (  
    df_without_duplicates  
    .select("brand", "model")  
    .show(10)  
    )
```

```
+-----+-----+  
| brand|          model|  
+-----+-----+  
|  Ford|Capri Extended Ra...|  
|Citroen|e-SpaceTourer XL ...|  
|Hyundai|IONIQ 6 Standard ...|  
|   NIO|      EL7 Long Range|  
|Porsche|      Macan Electric|  
|   BYD|SEAL U 87 kWh Design|  
|   Fiat|  500e Cabrio 42 kWh|  
|  Skoda|      Enyaq Coupe 60|  
|  Tesla|Model 3 Long Rang...|  
|   Audi|Q6 e-tron perform...|  
+-----+-----+  
only showing top 10 rows
```

---

## String & Column Operations

24. Add a new column `brand_model` by concatenating `brand` and `model` .

```
In [21]: df_with_brand_model = (  
    df  
    .withColumn("brand_model", psf.concat_ws(" ", "brand", "model"))  
    )
```

```
In [22]: (  
    df_with_brand_model  
    .select("brand", "model", "brand_model")  
    .show(10, truncate=False)  
    )
```

```

+-----+-----+-----+-----+
+-----+
|brand |model                               |brand_model
|
+-----+-----+-----+-----+
+-----+
|Abarth|500e Convertible                   |Abarth 500e Convertible
|
|Abarth|500e Hatchback                     |Abarth 500e Hatchback
|
|Abarth|600e Scorpionissima               |Abarth 600e Scorpionissima
|
|Abarth|600e Turismo                       |Abarth 600e Turismo
|
|Aiways|U5                                |Aiways U5
|
|Aiways|U6                                |Aiways U6
|
|Alfa   |Romeo Junior Elettrica 54 kWh          |Alfa Romeo Junior Elettrica 5
4 kWh
|
|Alfa   |Romeo Junior Elettrica 54 kWh Veloce|Alfa Romeo Junior Elettrica 5
4 kWh Veloce|
|Alpine|A290 Electric 180 hp                   |Alpine A290 Electric 180 hp
|
|Alpine|A290 Electric 220 hp                 |Alpine A290 Electric 220 hp
|
+-----+-----+-----+-----+
+-----+
only showing top 10 rows

```

---

25. Extract domain name from the `source_url` .

```

In [25]: (
    df
    .withColumn("domain", psf.regex_extract("source_url", r"https?://([^\s"]+\.?)"
    .select("source_url", "domain")
    .show(10, truncate=False)
)

```

```

+-----+
+-----+
|source_url
|domain      |
+-----+
+-----+
|https://ev-database.org/car/1904/Abarth-500e-Convertible
|ev-database.org|
|https://ev-database.org/car/1903/Abarth-500e-Hatchback
|ev-database.org|
|https://ev-database.org/car/3057/Abarth-600e-Scorpionissima
|ev-database.org|
|https://ev-database.org/car/3056/Abarth-600e-Turismo
|ev-database.org|
|https://ev-database.org/car/1678/Aiways-U5
|ev-database.org|
|https://ev-database.org/car/1766/Aiways-U6
|ev-database.org|
|https://ev-database.org/car/2184/Alfa-Romeo-Junior-Elettrica-54-kWh
|ev-database.org|
|https://ev-database.org/car/2185/Alfa-Romeo-Junior-Elettrica-54-kWh-Veloc
e|ev-database.org|
|https://ev-database.org/car/2268/Alpine-A290-Electric-180-hp
|ev-database.org|
|https://ev-database.org/car/2269/Alpine-A290-Electric-220-hp
|ev-database.org|
+-----+
+-----+
only showing top 10 rows

```

---

## Type Casting

26. Cast `battery_capacity_kWh` from float to integer.

```
In [26]: df_battery_in_int = (
          df.withColumn("battery_capacity_kWh_int", psf.col("battery_capacity_k
          .drop("battery_capacity_kWh")
          )

```

```
In [27]: df_battery_in_int.select("brand", "battery_capacity_kWh_int").show(10)

```

```

+-----+-----+
| brand|battery_capacity_kWh_int|
+-----+-----+
|Abarth|          37|
|Abarth|          37|
|Abarth|          50|
|Abarth|          50|
|Aiways|          60|
|Aiways|          60|
| Alfa|          50|
| Alfa|          50|
|Alpine|          52|
|Alpine|          52|
+-----+-----+
only showing top 10 rows

```

---

27. Convert `range_km` to string.

```

In [29]: df_range_km_in_string = (
    df
    .withColumn("range_km_str", psf.col("range_km").cast("string"))
    .drop("range_km")
)

```

```

In [32]: (
    df_range_km_in_string
    .select("brand", "range_km_str")
    .show(10)
)

```

```

+-----+-----+
| brand|range_km_str|
+-----+-----+
|Abarth|      225|
|Abarth|      225|
|Abarth|      280|
|Abarth|      280|
|Aiways|      315|
|Aiways|      350|
| Alfa|      320|
| Alfa|      310|
|Alpine|      310|
|Alpine|      305|
+-----+-----+
only showing top 10 rows

```

---

## Basic UDFs

28. Create a UDF to categorize cars as "High Range" (>300 km) or "Low Range".

```
In [15]: # Using pandas UDF
def categorize_cars(range_km: pd.Series) -> pd.Series:
    return range_km.apply(lambda x: "High Range" if x > 300 else "Low Ran
```

```
In [16]: categorize_udf = psf.pandas_udf(categorize_cars, returnType=pst.StringType)
```

```
In [17]: df_with_categories = (
    df
    .withColumn("Category", categorize_udf(psf.col("range_km")))
)
```

```
In [18]: (
    df_with_categories
    .select("brand", "range_km", "Category")
    .show(10)
)
```

```
+-----+-----+-----+
| brand|range_km|  Category|
+-----+-----+-----+
|Abarth|    225| Low Range|
|Abarth|    225| Low Range|
|Abarth|    280| Low Range|
|Abarth|    280| Low Range|
|Aiways|    315|High Range|
|Aiways|    350|High Range|
|  Alfa|    320|High Range|
|  Alfa|    310|High Range|
|Alpine|    310|High Range|
|Alpine|    305|High Range|
+-----+-----+-----+
only showing top 10 rows
```

---

29. Create a UDF to compute `power_density = battery_capacity_kWh / length_mm`.

```
In [20]: # Using pandas UDF
def compute_power_density(capacity: pd.Series, length: pd.Series) -> pd.S
    result = capacity * length
    return result
```

```
In [25]: compute_power_density_udf = psf.pandas_udf(compute_power_density, returnType=
```

```
In [26]: df_with_power_density = (
    df
    .withColumn("power_density", compute_power_density_udf(psf.col("batte
```

```
In [27]: (
    df_with_power_density
    .select("brand", "battery_capacity_kWh", "length_mm", "power_density")
    .show(10)
)
```

```

+-----+-----+-----+-----+
| brand|battery_capacity_kWh|length_mm|power_density|
+-----+-----+-----+-----+
|Abarth|          37.8|    3673|    138839.4|
|Abarth|          37.8|    3673|    138839.4|
|Abarth|          50.8|    4187|    212699.6|
|Abarth|          50.8|    4187|    212699.6|
|Aiways|          60.0|    4680|    280800.0|
|Aiways|          60.0|    4805|    288300.0|
| Alfa|          50.8|    4173|    211988.4|
| Alfa|          50.8|    4173|    211988.4|
|Alpine|          52.0|    3997|    207844.0|
|Alpine|          52.0|    3997|    207844.0|
+-----+-----+-----+-----+
only showing top 10 rows

```

---

## More Selections

30. Create a column `is_fast_charge_supported` where  
`fast_charging_power_kw_dc > 50`.

```

In [35]: df_with_new_col = (
          df
          .withColumn("is_fast_charge_supported", psf.expr("fast_charging_power

```

```

In [37]: (
          df_with_new_col
          .select("brand", "fast_charging_power_kw_dc", "is_fast_charge_support
          .show(10)
          )

```

```

+-----+-----+-----+-----+
| brand|fast_charging_power_kw_dc|is_fast_charge_supported|
+-----+-----+-----+-----+
|Abarth|          67|            true|
|Abarth|          67|            true|
|Abarth|          79|            true|
|Abarth|          79|            true|
|Aiways|          78|            true|
|Aiways|          78|            true|
| Alfa|          85|            true|
| Alfa|          85|            true|
|Alpine|          70|            true|
|Alpine|          70|            true|
+-----+-----+-----+-----+
only showing top 10 rows

```

---

31. Filter all rows where `car_body_type` is "SUV".

```

In [41]: df_SUV_only = (
          df

```

```
.where(psf.expr("car_body_type == 'SUV'))  
)
```

```
In [42]: (  
    df_SUV_only  
    .select("brand", "car_body_type")  
    .show(10)  
)
```

```
+-----+-----+  
| brand|car_body_type|  
+-----+-----+  
|Abarth|      SUV|  
|Abarth|      SUV|  
|Aiways|      SUV|  
|Aiways|      SUV|  
|  Alfa|      SUV|  
|  Alfa|      SUV|  
|  Audi|      SUV|  
|  Audi|      SUV|  
|  Audi|      SUV|  
|  Audi|      SUV|  
+-----+-----+  
only showing top 10 rows
```

---

## DataFrame Metadata

32. Show all column names and their data types.

```
In [45]: df.printSchema()
```

```
root  
|-- brand: string (nullable = true)  
|-- model: string (nullable = true)  
|-- top_speed_kmh: integer (nullable = true)  
|-- battery_capacity_kWh: double (nullable = true)  
|-- battery_type: string (nullable = true)  
|-- number_of_cells: integer (nullable = true)  
|-- torque_nm: integer (nullable = true)  
|-- efficiency_wh_per_km: integer (nullable = true)  
|-- range_km: integer (nullable = true)  
|-- acceleration_0_100_s: double (nullable = true)  
|-- fast_charging_power_kw_dc: integer (nullable = true)  
|-- fast_charge_port: string (nullable = true)  
|-- towing_capacity_kg: integer (nullable = true)  
|-- cargo_volume_l: integer (nullable = true)  
|-- seats: integer (nullable = true)  
|-- drivetrain: string (nullable = true)  
|-- segment: string (nullable = true)  
|-- length_mm: integer (nullable = true)  
|-- width_mm: integer (nullable = true)  
|-- height_mm: integer (nullable = true)  
|-- car_body_type: string (nullable = true)  
|-- source_url: string (nullable = true)
```

```
In [46]: df.dtypes
```

```
Out[46]: [('brand', 'string'),
          ('model', 'string'),
          ('top_speed_kmh', 'int'),
          ('battery_capacity_kWh', 'double'),
          ('battery_type', 'string'),
          ('number_of_cells', 'int'),
          ('torque_nm', 'int'),
          ('efficiency_wh_per_km', 'int'),
          ('range_km', 'int'),
          ('acceleration_0_100_s', 'double'),
          ('fast_charging_power_kw_dc', 'int'),
          ('fast_charge_port', 'string'),
          ('towing_capacity_kg', 'int'),
          ('cargo_volume_l', 'int'),
          ('seats', 'int'),
          ('drivetrain', 'string'),
          ('segment', 'string'),
          ('length_mm', 'int'),
          ('width_mm', 'int'),
          ('height_mm', 'int'),
          ('car_body_type', 'string'),
          ('source_url', 'string')]
```

---

33. Count total number of rows in the DataFrame.

```
In [47]: (
          df
          .count()
          )
```

```
Out[47]: 478
```

---

34. Check if all entries in the `drivetrain` column are the same.

```
In [48]: (
          df
          .select("drivetrain")
          .distinct()
          .show()
          )
```

```
+-----+
|drivetrain|
+-----+
|      FWD|
|      AWD|
|      RWD|
+-----+
```

---



35. Count the number of vehicles per number of seats.

```
In [50]: (
    df
    .groupBy("seats")
    .agg(psf.count("*").alias("count_per_seats"))
    .show()
)
```

```
+-----+-----+
|seats|count_per_seats|
+-----+-----+
|    6|              5|
|    5|            383|
|    9|             15|
|    4|             27|
|    8|              7|
|    7|             38|
|    2|              3|
+-----+-----+
```

---

## Medium Level (15 Questions)

Focus: Window functions, advanced aggregations, joins, UDFs, pivot/unpivot, JSON export, date functions (simulated), performance tuning basics.

### Window Functions & Advanced Aggregations

36. For each `brand`, rank vehicles by `range_km` in descending order.

```
In [54]: windowSpec = psf.Window.partitionBy("brand").orderBy(psf.col("range_km").
```

```
In [55]: df_with_range_rank = (
    df
    .withColumn("range_rank", psf.rank().over(windowSpec))
)
```

```
In [57]: (
    df_with_range_rank
    .select("brand", "range_km", "range_rank")
    .show(10)
)
```

```
+-----+-----+-----+
| brand|range_km|range_rank|
+-----+-----+-----+
|Abarth|    280|         1|
|Abarth|    280|         1|
|Abarth|    225|         3|
|Abarth|    225|         3|
|Aiways|    350|         1|
|Aiways|    315|         2|
| Alfa|    320|         1|
| Alfa|    310|         2|
|Alpine|    310|         1|
|Alpine|    305|         2|
+-----+-----+-----+
only showing top 10 rows
```

---

37. Calculate the average `battery_capacity_kWh` and standard deviation per `car_body_type` .

```
In [58]: df_with_avg_stdDev = (
          df
          .groupBy("car_body_type")
          .agg(psf.mean("battery_capacity_kWh").alias("average"),
              psf.stddev("battery_capacity_kWh").alias("standard deviation"))
          )
```

```
In [59]: df_with_avg_stdDev.show(10)
```

```
+-----+-----+-----+
|      car_body_type|      average|standard deviation|
+-----+-----+-----+
|      Hatchback|49.73684210526315| 14.51156720367469|
|      SUV|76.68729508196722|18.244932845760577|
|      Sedan| 86.8047619047619| 16.47529361010512|
| Liftback Sedan|85.52121212121212|12.055619995871156|
|Small Passenger Van|60.01489361702127|14.491211410389768|
|      Cabriolet|      58.08| 27.10289652417247|
|      Coupe|      92.5|13.435028842544403|
| Station/Estate|83.67407407407407|15.848309606356047|
+-----+-----+-----+
```

---

38. Find the vehicle with the longest `range_km` for each `segment` .

```
In [62]: (
          df
          .groupBy("segment")
          .agg(psf.max("range_km").alias("longest_range"))
          .show()
          )
```

segment	longest_range
JB - Compact	455
N - Passenger Van	370
A - Mini	225
B - Compact	330
F - Luxury	685
JD - Large	545
C - Medium	495
D - Large	565
E - Executive	555
JF - Luxury	540
I - Luxury	465
JA - Mini	300
JC - Medium	500
G - Sports	425
JE - Executive	610

## Pivoting

39. Pivot the data to show average `range_km` for each `car_body_type` per `drivetrain`.

```
In [76]: (
    df
    .groupBy("drivetrain")
    .pivot("car_body_type")
    .agg(psf.avg("range_km"))
    .transpose() # Just for printing it
    .show(10)
)
```

```

+-----+-----+-----+-----+
--+
|          key|          AWD|          FWD|          R
WD|
+-----+-----+-----+-----+
--+
|      Cabriolet|      395.0|      182.5|      42
5.0|
|      Coupe|      442.5|      NULL|      NU
LL|
|      Hatchback|      320.0|277.95454545454544|      39
2.5|
|      Liftback Sedan|      505.0|      NULL|      46
8.0|
|      SUV|424.1818181818182| 348.2258064516129| 395.8333333333333
33|
|      Sedan|      510.75|      362.5|511.470588235294
14|
|Small Passenger Van|      340.0|234.73684210526315|277.857142857142
83|
|      Station/Estate|      488.4375|      302.5|497.857142857142
83|
+-----+-----+-----+-----+
--+

```

## Advanced UDFs

40. Create a UDF to classify vehicles into: "City EV", "Highway EV", or "Performance EV" based on `acceleration_0_100_s` and `range_km`.

```

def classify_vehicle(accel, range_km): if accel is None or range_km is None: return "Unknown" if accel < 6: return
"Performance EV" elif range_km > 300: return "Highway EV" else: return "City EV" from pyspark.sql.functions import udf
from pyspark.sql.types import StringType classify_udf = udf(classify_vehicle, StringType()) df.withColumn("ev_type",
classify_udf("acceleration_0_100_s", "range_km")).show()

```

```

In [78]: # Using pandas UDFs for better performance
def classify_vehicle(acceleration: pd.Series, range_km: pd.Series) -> pd.
result = []
for a, r in zip(acceleration, range_km):
    if pd.isnull(a) or pd.isnull(r):
        result.append("Unknown")
    elif a < 6:
        result.append("Performance EV")
    elif r > 300:
        result.append("Highway EV")
    else:
        result.append("City EV")

return pd.Series(result)

```

```

In [79]: classify_vehicles_udf = psf.pandas_udf(classify_vehicle, returnType=pst.S

```

```

In [80]: df_with_classified_vehicles = (
df
.withColumn("class", classify_vehicles_udf(psf.col("acceleration_0_10
)

```

```
In [81]: (
    df_with_classified_vehicles
    .select("brand", "acceleration_0_100_s", "range_km", "class")
    .show(10)
)
```

brand	acceleration_0_100_s	range_km	class
Abarth	7.0	225	City EV
Abarth	7.0	225	City EV
Abarth	5.9	280	Performance EV
Abarth	6.2	280	City EV
Aiways	7.5	315	Highway EV
Aiways	7.0	350	Highway EV
Alfa	9.0	320	Highway EV
Alfa	6.0	310	Highway EV
Alpine	7.4	310	Highway EV
Alpine	6.4	305	Highway EV

only showing top 10 rows

## Joins (Synthetic example)

41. Assume you have another DataFrame with `brand` and `country`. Join it with the EV DataFrame on `brand`.

```
(
    df
    .join(df_country, on="brand", how="left")
    .show()
)
```

## Complex Filtering & Conditions

42. Find SUVs with a top speed over 180 km/h and efficiency under 160 Wh/km.

```
In [85]: df.columns
```

```
Out[85]: ['brand',
          'model',
          'top_speed_kmh',
          'battery_capacity_kWh',
          'battery_type',
          'number_of_cells',
          'torque_nm',
          'efficiency_wh_per_kmh',
          'range_kmh',
          'acceleration_0_100_s',
          'fast_charging_power_kw_dc',
          'fast_charge_port',
          'towing_capacity_kg',
          'cargo_volume_l',
          'seats',
          'drivetrain',
          'segment',
          'length_mm',
          'width_mm',
          'height_mm',
          'car_body_type',
          'source_url']
```

```
In [87]: (
    df
    .select("brand", "model", "top_speed_kmh", "efficiency_wh_per_kmh")
    .filter(psf.col("car_body_type") == "SUV")
    .filter(psf.col("top_speed_kmh") > 180)
    .filter(psf.col("efficiency_wh_per_kmh") < 160)
    .show(10)
)
```

brand	model	top_speed_kmh	efficiency_wh_per_kmh
Abarth	600e Scorpionissima	200	158
Abarth	600e Turismo	200	158
Audi	Q6 e-tron Sportback	210	139
Audi	Q6 e-tron Sportba...	210	145
Audi	Q6 e-tron Sportba...	210	149
DS	N°8 AWD Long Range	190	146
DS	N°8 FWD	190	141
DS	N°8 FWD Long Range	190	136
Genesis	GV60 Premium	185	143
Genesis	GV60 Sport	200	157

only showing top 10 rows

## JSON Export & Schema

43. Write the DataFrame to a JSON file with inferred schema.

```
In [92]: file_path = "exported_json"
```

```
In [93]: (
    df
    .write
    .mode("overwrite")
    .format("json")
    .save(file_path)
)
```

---

44. Infer schema manually using `StructType` and load the json using it.

```
In [94]: file_path = "exported_json/part-00000-e1c2171c-df5c-4740-96bd-089ebb96ea1"
```

```
In [97]: schema = pst.StructType([
    pst.StructField("brand", pst.StringType(), True),
    pst.StructField("model", pst.StringType(), True),
    pst.StructField("top_speed_kmh", pst.IntegerType(), True),
    pst.StructField("battery_capacity_kWh", pst.DoubleType(), True),
    pst.StructField("battery_type", pst.StringType(), True),
    pst.StructField("number_of_cells", pst.DoubleType(), True),
    pst.StructField("torque_nm", pst.DoubleType(), True),
    pst.StructField("efficiency_wh_per_km", pst.IntegerType(), True),
    pst.StructField("range_km", pst.IntegerType(), True),
    pst.StructField("acceleration_0_100_s", pst.DoubleType(), True),
    pst.StructField("fast_charging_power_kw_dc", pst.DoubleType(), True),
    pst.StructField("fast_charge_port", pst.StringType(), True),
    pst.StructField("towing_capacity_kg", pst.DoubleType(), True),
    pst.StructField("cargo_volume_l", pst.IntegerType(), True),
    pst.StructField("seats", pst.IntegerType(), True),
    pst.StructField("drivetrain", pst.StringType(), True),
    pst.StructField("segment", pst.StringType(), True),
    pst.StructField("length_mm", pst.IntegerType(), True),
    pst.StructField("width_mm", pst.IntegerType(), True),
    pst.StructField("height_mm", pst.IntegerType(), True),
    pst.StructField("car_body_type", pst.StringType(), True),
    pst.StructField("source_url", pst.StringType(), True)
])
```

```
In [99]: df_read_from_json = (
    spark
    .read
    .schema(schema)
    .json(file_path)
)
```

```
In [100]: (
    df_read_from_json
    .select("brand", "model", "car_body_type", "source_url")
    .show(10)
)
```

brand	model	car_body_type	source_url
Abarth	500e Convertible	Hatchback	https://ev-databa...
Abarth	500e Hatchback	Hatchback	https://ev-databa...
Abarth	600e Scorpionissima	SUV	https://ev-databa...
Abarth	600e Turismo	SUV	https://ev-databa...
Aiways	U5	SUV	https://ev-databa...
Aiways	U6	SUV	https://ev-databa...
Alfa Romeo	Junior Elet...	SUV	https://ev-databa...
Alfa Romeo	Junior Elet...	SUV	https://ev-databa...
Alpine	A290 Electric 180 hp	Hatchback	https://ev-databa...
Alpine	A290 Electric 220 hp	Hatchback	https://ev-databa...

only showing top 10 rows

## Unpivot (Melt-like)

45. Transform the column-based format into a long format for `top_speed_kmh`, `range_km`, and `efficiency_wh_per_km`.

```
In [103... from pyspark.sql.functions import posexplode, array, struct, lit, explode

unpivoted = df.select(
    "brand", "model",
    explode(array(
        struct(lit("top_speed_kmh").alias("metric"), col("top_speed_kmh")),
        struct(lit("range_km").alias("metric"), col("range_km").alias("value")),
        struct(lit("efficiency_wh_per_km").alias("metric"), col("efficiency_wh_per_km").alias("value"))
    )).alias("kv"))
    ).select("brand", "model", col("kv.metric"), col("kv.value"))

unpivoted.show()
```



brand	model	metric	value
Abarth	500e Convertible	top_speed_kmh	155
Abarth	500e Convertible	range_km	225
Abarth	500e Convertible	efficiency_wh_per_km	156
Abarth	500e Hatchback	top_speed_kmh	155
Abarth	500e Hatchback	range_km	225
Abarth	500e Hatchback	efficiency_wh_per_km	149
Abarth	600e Scorpionissima	top_speed_kmh	200
Abarth	600e Scorpionissima	range_km	280
Abarth	600e Scorpionissima	efficiency_wh_per_km	158
Abarth	600e Turismo	top_speed_kmh	200
Abarth	600e Turismo	range_km	280
Abarth	600e Turismo	efficiency_wh_per_km	158
Aiways	U5	top_speed_kmh	150
Aiways	U5	range_km	315
Aiways	U5	efficiency_wh_per_km	156
Aiways	U6	top_speed_kmh	160
Aiways	U6	range_km	350
Aiways	U6	efficiency_wh_per_km	150
Alfa Romeo	Junior Elet...	top_speed_kmh	150
Alfa Romeo	Junior Elet...	range_km	320

only showing top 20 rows

below is the efficient version (using pandas UDF)

```

from pyspark.sql.functions import pandas_udf
from pyspark.sql.types import StructType, StructField,
StringType, DoubleType
import pandas as pd

# Define output schema for the long format
output_schema = StructType([
    StructField("brand", StringType()),
    StructField("model", StringType()),
    StructField("metric", StringType()),
    StructField("value", DoubleType())
])

# Pandas UDF for unpivoting
@pandas_udf(output_schema)
def unpivot_udf(pdf: pd.DataFrame) -> pd.DataFrame:
    id_vars = ["brand", "model"]
    value_vars = ["top_speed_kmh", "range_km",
"efficiency_wh_per_km"]

    # Melt to long format
    melted = pd.melt(
        pdf,
        id_vars=id_vars,
        value_vars=value_vars,
        var_name="metric",
        value_name="value"
    )
    return melted

```

```
unpivoted = df.groupBy("brand", "model").apply(unpivot_udf)
```

---

## Map & ReduceByKey (RDD-style)

46. Convert DataFrame to RDD and find average `battery_capacity_kWh` per brand using `reduceByKey`.

```
rdd = df.select("brand", "battery_capacity_kWh").rdd \
    .filter(lambda x: x[1] is not None) \
    .map(lambda x: (x[0], (x[1], 1)))

brand_avg = rdd.reduceByKey(lambda a, b: (a[0] + b[0], a[1] +
b[1])) \
    .mapValues(lambda x: x[0] / x[1])

brand_avg.collect()
```

---

## Performance Optimization

47. Cache the DataFrame and perform a `groupBy` operation. Measure execution time before and after caching.

```
In [104... import time

start = time.time()
df.groupBy("car_body_type").count().show()
print("Without cache:", time.time() - start)

df.cache()

start = time.time()
df.groupBy("car_body_type").count().show()
print("With cache:", time.time() - start)
```

car_body_type	count
Hatchback	57
SUV	244
Sedan	63
Liftback Sedan	33
Small Passenger Van	47
Cabriolet	5
Coupe	2
Station/Estate	27

Without cache: 0.4533965587615967

car_body_type	count
Hatchback	57
SUV	244
Sedan	63
Liftback Sedan	33
Small Passenger Van	47
Cabriolet	5
Coupe	2
Station/Estate	27

With cache: 0.3217630386352539

---

## Nested Column Creation

48. Create a struct column called `performance` with `top_speed_kmh`, `acceleration_0_100_s`, and `torque_nm`.

```
df.withColumn("performance", struct("top_speed_kmh", "torque_nm",
"acceleration_0_100_s")).select("brand", "model",
"performance").show(truncate=False)
```

```
In [105... df_with_performance_col = (
    df
    .withColumn("performance", psf.struct("top_speed_kmh", "torque_nm", "
    .select("brand", "model", "performance")
)
```

```
In [106... df_with_performance_col.show(10)
```

```
+-----+-----+-----+
| brand|          model|   performance|
+-----+-----+-----+
|Abarth|    500e Convertible|{155, 235, 7.0}|
|Abarth|    500e Hatchback|{155, 235, 7.0}|
|Abarth| 600e Scorpionissima|{200, 345, 5.9}|
|Abarth|    600e Turismo|{200, 345, 6.2}|
|Aiways|              U5|{150, 310, 7.5}|
|Aiways|              U6|{160, 315, 7.0}|
| Alfa|Romeo Junior Elet...|{150, 260, 9.0}|
| Alfa|Romeo Junior Elet...|{200, 345, 6.0}|
|Alpine|A290 Electric 180 hp|{160, 285, 7.4}|
|Alpine|A290 Electric 220 hp|{170, 300, 6.4}|
+-----+-----+-----+
only showing top 10 rows
```

---

## Exploding Arrays

49. Create an array column of all performance metrics and explode it row-wise.

```
In [107... (
    df
    .withColumn("performance_array", array("top_speed_kmh", "torque_nm",
    .withColumn("exploded_metric", explode("performance_array"))
    .select("brand", "model", "exploded_metric")
    .show(10)
)
```

```
+-----+-----+-----+
| brand|          model|exploded_metric|
+-----+-----+-----+
|Abarth|    500e Convertible|      155.0|
|Abarth|    500e Convertible|      235.0|
|Abarth|    500e Convertible|       7.0|
|Abarth|    500e Hatchback|      155.0|
|Abarth|    500e Hatchback|      235.0|
|Abarth|    500e Hatchback|       7.0|
|Abarth| 600e Scorpionissima|      200.0|
|Abarth| 600e Scorpionissima|      345.0|
|Abarth| 600e Scorpionissima|       5.9|
|Abarth|    600e Turismo|      200.0|
+-----+-----+-----+
only showing top 10 rows
```

---

## Complex Sorting

50. Sort cars by acceleration (ascending), then by torque (descending), and then by range.

```
In [109... (
    df
    .orderBy(
        psf.col("acceleration_0_100_s").asc(),
```

```

    psf.col("torque_nm").desc(),
    psf.col("range_km").desc()
)
.select("brand", "model", "acceleration_0_100_s", "torque_nm", "range_km")
.show()
)

```

brand	model	acceleration_0_100_s	torque_nm	range_km
Porsche	Taycan Turbo GT W...	2.2	1340	475
Porsche	Taycan Turbo GT	2.3	1340	475
Tesla	Model S Plaid	2.3	NULL	560
Porsche	Taycan Turbo S	2.4	1110	525
Porsche	Taycan Turbo S Sp...	2.4	1110	505
Porsche	Taycan Turbo S Cr...	2.5	1110	485
Audi	e-tron GT RS perf...	2.5	1027	525
Maserati	GranTurismo Folgore	2.7	1350	420
Porsche	Taycan Turbo	2.7	940	535
Porsche	Taycan Turbo Spor...	2.7	940	505
Tesla	Model X Plaid	2.7	NULL	465
Maserati	GranCabrio Folgore	2.8	1350	395
Lotus	Emeya R	2.8	985	465
Porsche	Taycan Turbo Cros...	2.8	940	495
Audi	e-tron GT RS	2.8	865	525
Lotus	Eletre R	2.9	985	455
Lucid	Air Grand Touring	3.0	1200	665
Tesla	Model 3 Performan...	3.2	741	490
MG	Cyberster GT	3.2	725	395
Lucid	Air Touring	3.2	NULL	580

only showing top 20 rows